skills_manifold

April 23, 2026

1 Skill Manifold — Comparison B (Gromov–Wasserstein)

Interactive walkthrough of `pipeline/skill_manifold_gw.py`. We compare two independent **skill manifolds**:

1. **JIGSAWS** — dVRK robotic surgery video: three tasks (Suturing, Knot_Tying, Needle_Passing), 8 surgeons, per-trial GRS + 6 OSATS sub-scores.
2. **EEG/Eye simulator** — da Vinci Mimic: 25 participants × 27 tasks × per-try performance 0–100.

The datasets have different dimensionality, different subjects, and no shared task taxonomy. **Gromov–Wasserstein** is the right tool because it compares intra-set pairwise geometry only.

Question. After conditioning out task, subject, and other nuisances, does the residual *skill geometry* on each side have the same structure?

Run all cells top-to-bottom; expect 30 s on the real dataset.

[1]:
```
%load_ext autoreload
%autoreload 2

import sys, json
from pathlib import Path

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

REPO = Path.cwd().resolve()
if REPO.name != 'surgical_gestures_skills_manifold':
    REPO = REPO.parent   # notebook is in pipeline/
sys.path.insert(0, str(REPO / 'src'))

from skill_manifold.io import (
    CONFIG_DIR, JIGSAWS_OSATS_COLUMNS, data_root, load_config,
)
from skill_manifold.features_jigsaws import (
    build_jigsaws_feature_frame,
    feature_column_names as jigsaws_feat_cols,
    zscore_features,
```

```python
)
from skill_manifold.features_eeg_eye import (
    build_eeg_eye_feature_frame,
    feature_column_names as eeg_feat_cols,
    load_task_module_map,
)
from skill_manifold.residualize import residualize
from skill_manifold.binning import TIER_NAMES, add_tier_column
from skill_manifold.rdms import centroid_rdm, pairwise_cosine_rdm
from skill_manifold.gw import (
    entropic_gromov_wasserstein,
    gromov_wasserstein_centroids,
    permutation_null_centroid,
)

cfg = load_config('skill_manifold.yaml')
ROOT = data_root(None)
print('data root:', ROOT)
print('config:   ', cfg)
```

```
---------------------------------------------------------------------------
AttributeError                            Traceback (most recent call last)
AttributeError: 'MessageFactory' object has no attribute 'GetPrototype'
```

```
---------------------------------------------------------------------------
AttributeError                            Traceback (most recent call last)
AttributeError: 'MessageFactory' object has no attribute 'GetPrototype'
```

```
---------------------------------------------------------------------------
AttributeError                            Traceback (most recent call last)
AttributeError: 'MessageFactory' object has no attribute 'GetPrototype'
```

```
---------------------------------------------------------------------------
AttributeError                            Traceback (most recent call last)
AttributeError: 'MessageFactory' object has no attribute 'GetPrototype'
```

```
---------------------------------------------------------------------------
AttributeError                            Traceback (most recent call last)
AttributeError: 'MessageFactory' object has no attribute 'GetPrototype'
```

```
---------------------------------------------------------------------------
AttributeError                            Traceback (most recent call last)
```

```
AttributeError: 'MessageFactory' object has no attribute 'GetPrototype'
```

```
data root: /Users/michaelhaidar/Library/Mobile
Documents/com~apple~CloudDocs/Documents/Vanderbilt/Fall_25/Surgical
Robotics/Surgical_Gestures
config:    {'seed': 1337, 'n_perms': 1000, 'gw_epsilon': 0.01,
'subsample_per_tier': 100, 'trial_null_n_subsamples': 30, 'gripper_open_eps':
1e-06, 'economy_clip': '1.0e6', 'try_filter': 1, 'gestures_pool': ['G1', 'G2',
'G3', 'G4', 'G5', 'G6', 'G8', 'G9', 'G10', 'G11', 'G12', 'G13', 'G14', 'G15'],
'osats_axes': ['respect_for_tissue', 'suture_needle_handling',
'time_and_motion', 'flow_of_operation', 'overall_performance',
'quality_of_final_product']}
```

1.1 Step 1 — JIGSAWS feature extraction

Per trial: 14-dim gesture histogram + per-arm Slave-Left/Slave-Right kinematics summaries (12 scalars each) + trial duration → **39 dims**. Z-scored across the pooled JIGSAWS set.

[2]:
```
gp = list(cfg['gestures_pool'])
jf = build_jigsaws_feature_frame(
    ROOT, gp,
    gripper_open_eps=float(cfg['gripper_open_eps']),
    economy_clip=float(cfg['economy_clip']),
)
j_cols = jigsaws_feat_cols(gp)
jf = zscore_features(jf, j_cols)
print(f'JIGSAWS trials: {len(jf)}   feature dim: {len(j_cols)}')
jf[['trial_id', 'task', 'surgeon', 'skill', 'grs_total'] + j_cols[:4]].head()
```

```
JIGSAWS trials: 103   feature dim: 39
```

[2]:
```
          trial_id        task surgeon skill  grs_total   gest_G1   gest_G2  \
0  Knot_Tying_B001  Knot_Tying       B     N         13 -1.016041 -0.948707
1  Knot_Tying_B002  Knot_Tying       B     N          9 -1.016041 -0.948707
2  Knot_Tying_B003  Knot_Tying       B     N         11 -1.016041 -0.948707
3  Knot_Tying_B004  Knot_Tying       B     N          8 -1.016041 -0.948707
4  Knot_Tying_C001  Knot_Tying       C     I         20 -1.016041 -0.948707

    gest_G3   gest_G4
0 -1.077075 -1.233897
1 -1.077075 -1.233897
2 -1.077075 -1.233897
3 -1.077075 -1.233897
4 -1.077075 -1.233897
```

[3]:
```
fig, axes = plt.subplots(1, 2, figsize=(10, 3.5))
jf['grs_total'].hist(bins=20, ax=axes[0], color='#4575b4', edgecolor='white')
axes[0].set_title('JIGSAWS GRS total (pooled)')
```

```
axes[0].set_xlabel('GRS total'); axes[0].set_ylabel('count')
jf['task'].value_counts().plot.bar(ax=axes[1], color='#984ea3')
axes[1].set_title('JIGSAWS trials per task'); axes[1].tick_params(axis='x',␣
↪rotation=0)
plt.tight_layout(); plt.show()
```

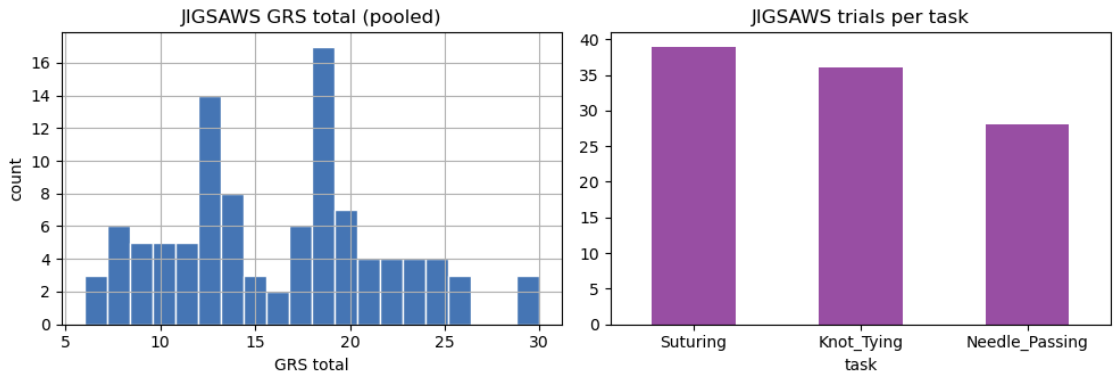

1.2 Step 2 — EEG/Eye simulator feature extraction

Per Try==1 trial: mean-pooled Phase 1 `baseline_embeddings` (64) + mean-pooled `pc_embeddings` (64) + 18-dim eye summary reconstructed from Phase 2 components → **146 dims**.

> The cached `fingerprint` key is absent on ~91% of trials, so we rebuild an equivalent vector from `eye_state_summary`, `eye_transition_summary`, `pupil_summary`, `event_summary` — all of which are present in every trial.

[4]:
```
mod_map = load_task_module_map(CONFIG_DIR / 'skill_manifold_task_modules.yaml')
ef = build_eeg_eye_feature_frame(ROOT, mod_map,␣
↪try_filter=int(cfg['try_filter']))
e_cols = eeg_feat_cols()
# Z-score per feature on the EEG/Eye side.
for c in e_cols:
    v = ef[c].to_numpy(float); sd = v.std()
    ef[c] = ((v - v.mean()) / sd) if sd > 1e-12 else 0.0
print(f'EEG/Eye trials (Try==1, both cache hits): {len(ef)}   feature dim:␣
↪{len(e_cols)}')
ef[['trial_id', 'subject_id', 'task_id', 'task_module', 'age', 'performance']].
↪head()
```

```
EEG/Eye trials (Try==1, both cache hits): 510   feature dim: 146
```

[4]:

	trial_id	subject_id	task_id	task_module	age	performance
0	10_10_1	10	10	camera_scaling	34.0	95.0
1	10_11_1	10	11	camera_scaling	34.0	63.0
2	10_12_1	10	12	ring_walk	34.0	97.0

```
3  10_13_1          10        13        ring_walk  34.0          79.0
4  10_14_1          10        14        ring_walk  34.0          82.0
```

[5]:
```
fig, axes = plt.subplots(1, 2, figsize=(10, 3.5))
ef['performance'].hist(bins=20, ax=axes[0], color='#1a9850', edgecolor='white')
axes[0].set_title('EEG/Eye Mimic performance (out of 100)')
axes[0].set_xlabel('performance'); axes[0].set_ylabel('count')
ef['task_module'].value_counts().plot.bar(ax=axes[1], color='#d73027')
axes[1].set_title('EEG/Eye trials per task-module')
axes[1].tick_params(axis='x', rotation=30); axes[1].set_xlabel('')
plt.tight_layout(); plt.show()
```

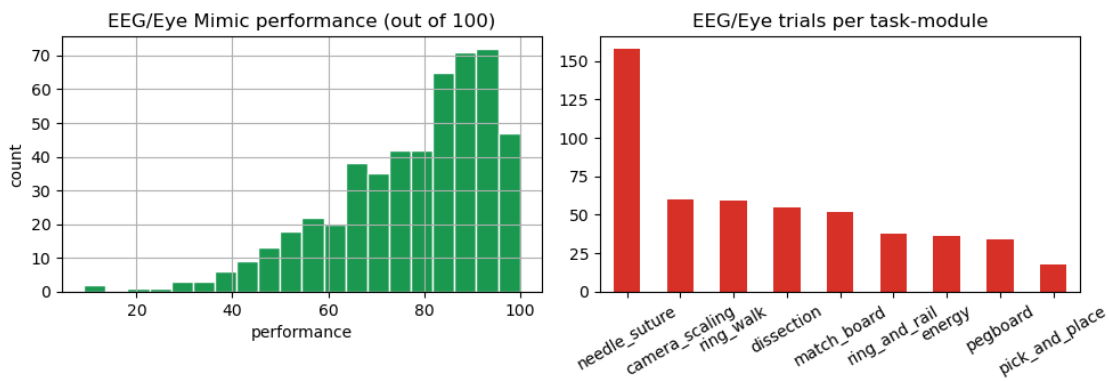

1.3 Step 3 — Nuisance residualization

We fit `feature ~ nuisances` with OLS and keep the residuals. **JIGSAWS nuisances:** task (3 levels), surgeon (8 levels), trial index within (surgeon,task). **EEG/Eye nuisances:** task_module (9 levels), subject_id (25 levels), dominant_hand (2 levels), age (ordinal).

[6]:
```
jr = residualize(jf, j_cols,
                 categorical=['task', 'surgeon'],
                 ordinal=['trial_index_within_surgeon_task'])
er = residualize(ef, e_cols,
                 categorical=['task_module', 'subject_id', 'dominant_hand'],
                 ordinal=['age'])
print(f'JIGSAWS  max R^2 on features: {jr.r2_per_feature.max():.3f}   '
      f'max R^2 after residualizing: {jr.post_fit_r2.max():.2e}')
print(f'EEG/Eye  max R^2 on features: {er.r2_per_feature.max():.3f}   '
      f'max R^2 after residualizing: {er.post_fit_r2.max():.2e}')

fig, axes = plt.subplots(1, 2, figsize=(10, 3.5))
axes[0].hist(jr.r2_per_feature, bins=25, color='#4575b4', edgecolor='white')
axes[0].set_title('JIGSAWS: nuisance R² per feature')
axes[0].set_xlabel('R²'); axes[0].set_ylabel('n features')
```

```
axes[1].hist(er.r2_per_feature, bins=25, color='#1a9850', edgecolor='white')
axes[1].set_title('EEG/Eye: nuisance R² per feature')
axes[1].set_xlabel('R²'); axes[1].set_ylabel('n features')
plt.tight_layout(); plt.show()
```

```
JIGSAWS  max R^2 on features: 0.945   max R^2 after residualizing: 2.22e-16
EEG/Eye  max R^2 on features: 0.655   max R^2 after residualizing: 3.33e-16
```

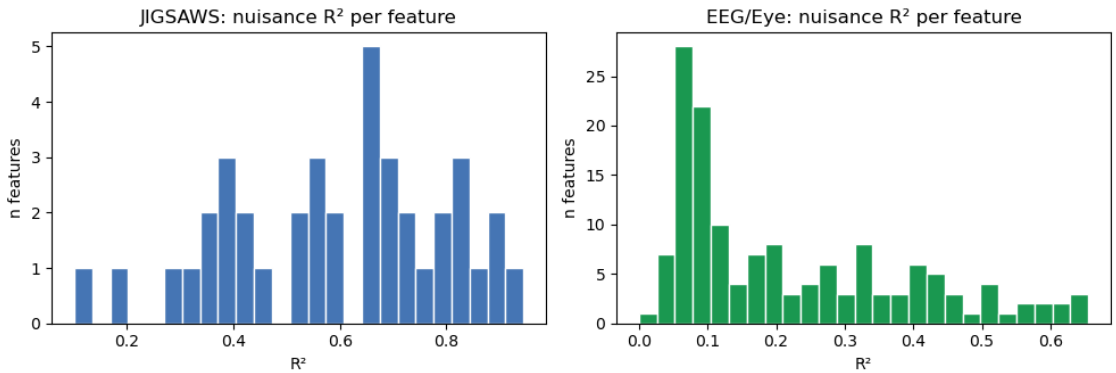

1.4 Step 4 — Tertile binning

JIGSAWS → tertiles of GRS total. EEG/Eye → tertiles of Mimic performance. These are the **Low / Mid / High** skill tiers used everywhere downstream.

[7]:
```
jr_df, j_cut = add_tier_column(jr.residuals, 'grs_total')
er_df, e_cut = add_tier_column(er.residuals, 'performance')
print(f'JIGSAWS GRS cutoffs: q33 = {j_cut.q33:g}, q66 = {j_cut.q66:g}')
print(f'EEG/Eye performance cutoffs: q33 = {e_cut.q33:g}, q66 = {e_cut.q66:g}')

cts = pd.DataFrame({
    'JIGSAWS': jr_df['tier'].value_counts().reindex(TIER_NAMES),
    'EEG/Eye': er_df['tier'].value_counts().reindex(TIER_NAMES),
})
display(cts)

ax = cts.plot.bar(figsize=(6, 3.5), color=['#4575b4', '#1a9850'])
ax.set_title('Trials per tier'); ax.set_xlabel('tier'); ax.
 ↪tick_params(axis='x', rotation=0)
plt.tight_layout(); plt.show()
```

```
JIGSAWS GRS cutoffs: q33 = 13, q66 = 19
EEG/Eye performance cutoffs: q33 = 72, q66 = 88
```

	JIGSAWS	EEG/Eye
tier		
Low	38	171

```
Mid         36       180
High        29       159
```

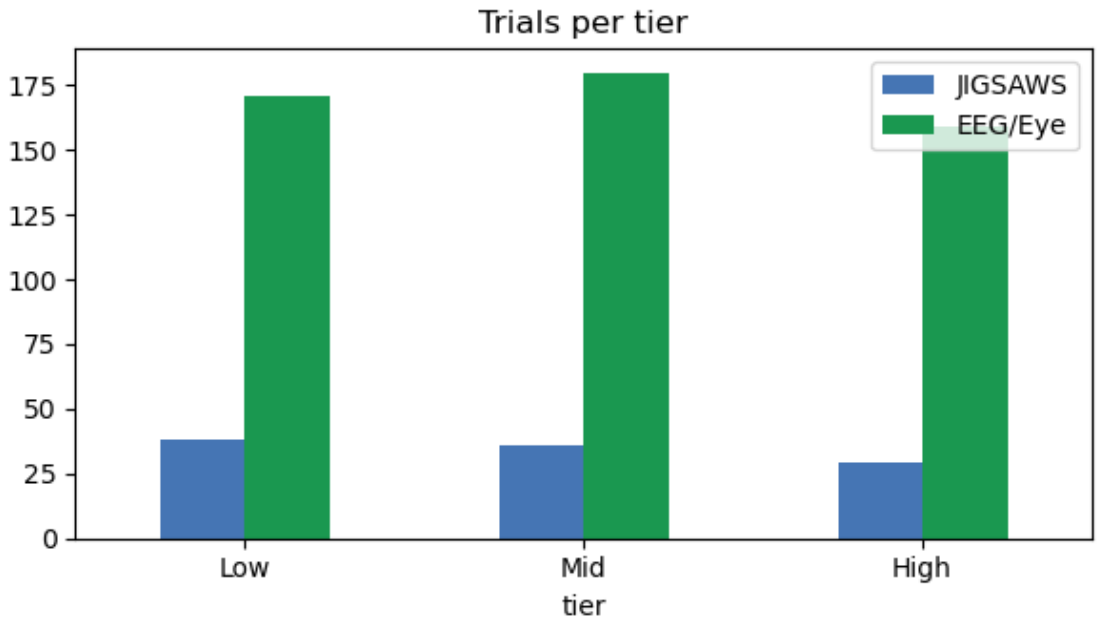

1.5 Step 5 — Centroid RDMs (Low / Mid / High)

Within each dataset: mean residualized feature vector per tier → 3 centroids → 3×3 cosine-distance RDM.

These are the two things GW compares.

[8]:
```
jf_mat = jr_df[j_cols].to_numpy(float)
ef_mat = er_df[e_cols].to_numpy(float)
RDM_J = centroid_rdm(jf_mat, jr_df['tier'].to_numpy(), TIER_NAMES)
RDM_E = centroid_rdm(ef_mat, er_df['tier'].to_numpy(), TIER_NAMES)

def _heat(ax, M, title):
    im = ax.imshow(M, cmap='viridis')
    ax.set_xticks(range(3), TIER_NAMES); ax.set_yticks(range(3), TIER_NAMES)
    for i in range(3):
        for j in range(3):
            ax.text(j, i, f'{M[i,j]:.3f}', ha='center', va='center',
                    color='white' if M[i,j] < M.max() * 0.6 else 'black',␣
 ↪fontsize=9)
    ax.set_title(title)
    return im

fig, axes = plt.subplots(1, 2, figsize=(10, 4))
```

```
im1 = _heat(axes[0], RDM_J, 'JIGSAWS centroid RDM')
im2 = _heat(axes[1], RDM_E, 'EEG/Eye centroid RDM')
fig.colorbar(im1, ax=axes[0], fraction=0.046)
fig.colorbar(im2, ax=axes[1], fraction=0.046)
plt.tight_layout(); plt.show()
```

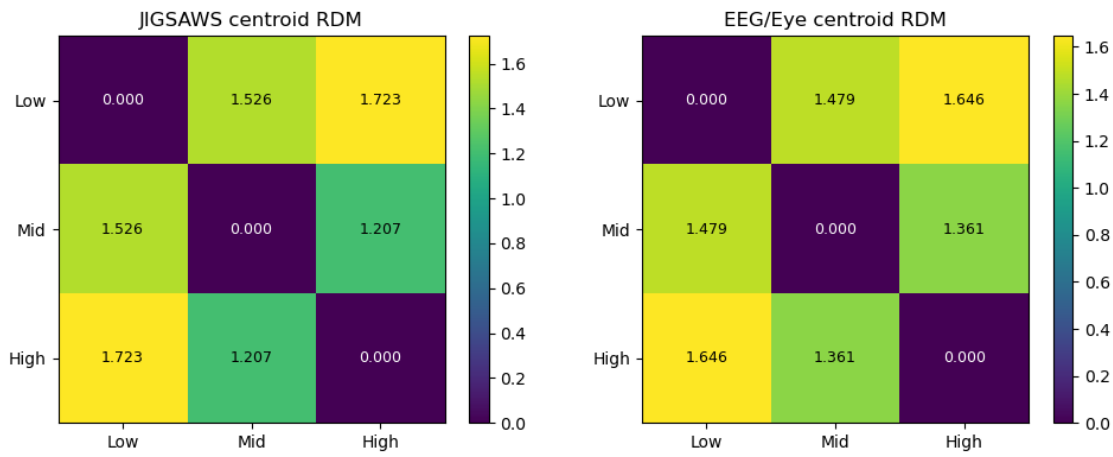

1.6 Step 6 — Gromov–Wasserstein distance + coupling

With uniform marginals and square loss, GW finds the transport plan that best aligns the two intra-set geometries. The **coupling** matrix tells us which source tier is mapped to which target tier.

[9]:
```
gw = gromov_wasserstein_centroids(RDM_J, RDM_E, TIER_NAMES)
print(f'GW distance = {gw.distance:.6f}')
print(f'Argmax tier-to-tier assignment: {gw.argmax_assignment}')

fig, ax = plt.subplots(figsize=(4.5, 4))
im = ax.imshow(gw.coupling, cmap='magma')
ax.set_xticks(range(3), TIER_NAMES); ax.set_yticks(range(3), TIER_NAMES)
ax.set_xlabel('EEG/Eye'); ax.set_ylabel('JIGSAWS')
for i in range(3):
    for j in range(3):
        ax.text(j, i, f'{gw.coupling[i,j]:.3f}', ha='center', va='center',
                color='white' if gw.coupling[i,j] < 0.2 else 'black',␣
 ↪fontsize=9)
ax.set_title(f'GW coupling (distance = {gw.distance:.4f})')
fig.colorbar(im, ax=ax, fraction=0.046)
plt.tight_layout(); plt.show()
```

```
GW distance = 0.007089
Argmax tier-to-tier assignment: {'Low': 'Low', 'Mid': 'Mid', 'High': 'High'}
```

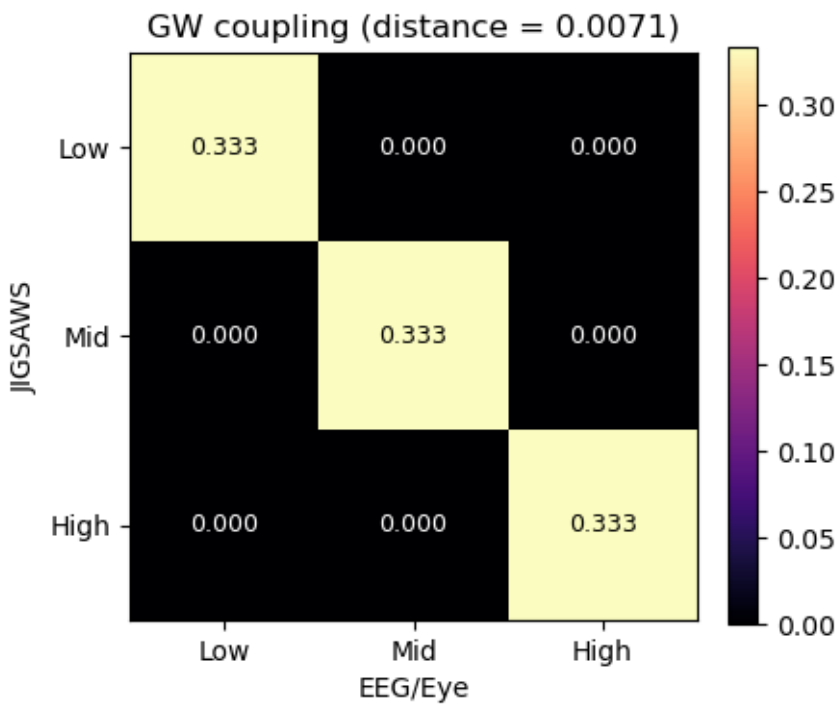

1.7 Step 7 — Permutation null

Under the null, tier labels carry no feature-space structure. We shuffle tier assignments on each side independently (preserving tier counts), recompute centroids → RDMs → GW, and compare the observed value to the null distribution.

p = (1 + #{null observed}) / (1 + n). **Small p means the observed alignment is unusually tight.**

[10]:
```
null = permutation_null_centroid(
    jf_mat, jr_df['tier'].to_numpy(),
    ef_mat, er_df['tier'].to_numpy(),
    TIER_NAMES,
    n_perms=int(cfg['n_perms']),
    seed=int(cfg['seed']),
)
print(f"observed GW     = {null['observed']:.6f}")
print(f"null mean ± std = {null['null_mean']:.6f} ± {null['null_std']:.6f}")
print(f"p-value         = {null['p_value']:.4f}")
print(f"z-score         = {null['z_score']:.3f}")
print(f"n permutations  = {null['n_permutations']}  (degenerate dropped:␣
↪{null['n_degenerate']})")

fig, ax = plt.subplots(figsize=(7, 3.8))
```

```
ax.hist(null['null'], bins=40, color='#999', edgecolor='white')
ax.axvline(null['observed'], color='red', linestyle='--',
           label=f"observed = {null['observed']:.4f}")
ax.set_xlabel('GW distance under tier-shuffle null')
ax.set_ylabel('count')
ax.set_title(f"Permutation null (p={null['p_value']:.3f}, z={null['z_score']:.
 ↪2f})")
ax.legend()
plt.tight_layout(); plt.show()
```

```
observed GW     = 0.007089
null mean ± std = 0.011142 ± 0.015636
p-value         = 0.5714
z-score         = -0.259
n permutations  = 1000  (degenerate dropped: 0)
```

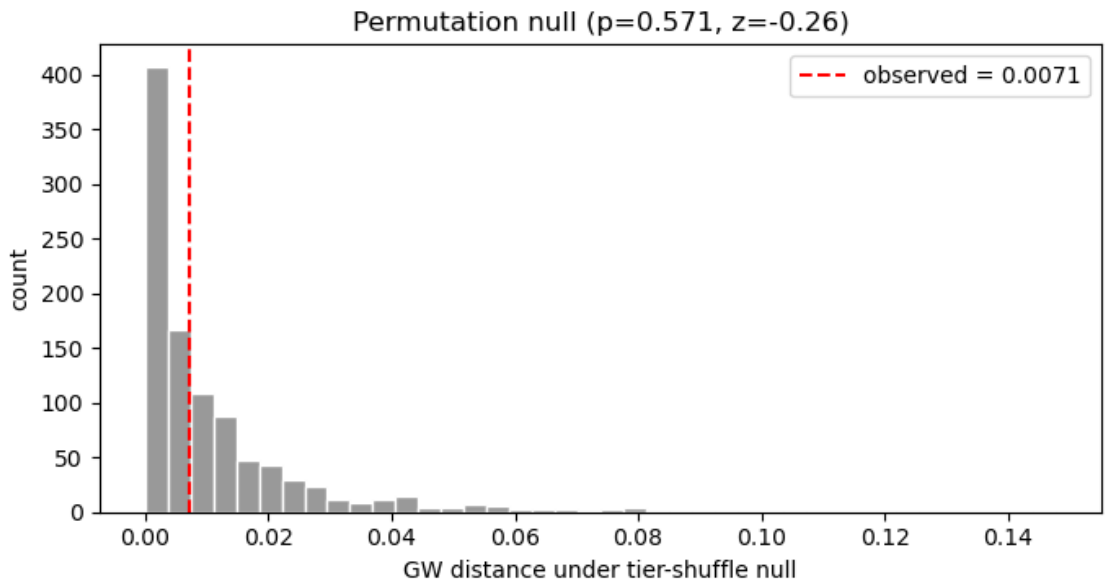

1.8 Step 8 — OSATS axis-alignment

The JIGSAWS meta files record six OSATS sub-scores per trial. For each axis we re-bin JIGSAWS trials into tertiles of that single axis, rebuild `RDM_J_axis`, and run GW against the same `RDM_E`.

Which sub-dimension of expert-rated skill aligns tightest with the Mimic composite?

[11]:
```
osats_rows = []
null_cis = []
for axis in JIGSAWS_OSATS_COLUMNS:
    colname = f'osats_{axis}'
    sub, _cut = add_tier_column(jr_df, colname, tier_col=f'tier_{axis}')
```

```
    R = centroid_rdm(sub[j_cols].to_numpy(float),
                     sub[f'tier_{axis}'].to_numpy(), TIER_NAMES)
    g = gromov_wasserstein_centroids(R, RDM_E, TIER_NAMES)
    n = permutation_null_centroid(
        sub[j_cols].to_numpy(float), sub[f'tier_{axis}'].to_numpy(),
        ef_mat, er_df['tier'].to_numpy(),
        TIER_NAMES, n_perms=max(200, int(cfg['n_perms']) // 5),
        seed=int(cfg['seed']))
    osats_rows.append({
        'axis': axis, 'gw': g.distance, 'p': n['p_value'], 'z': n['z_score'],
    })
    null_cis.append((float(np.quantile(n['null'], 0.025)),
                     float(np.quantile(n['null'], 0.975))))
osats_df = pd.DataFrame(osats_rows).set_index('axis')
display(osats_df.style.format({'gw': '{:.4f}', 'p': '{:.4f}', 'z': '{:+.2f}'}))

fig, ax = plt.subplots(figsize=(9, 3.8))
x = np.arange(len(osats_df))
ax.bar(x, osats_df['gw'], color='#4575b4', label='observed GW')
lo = np.array([c[0] for c in null_cis]); hi = np.array([c[1] for c in null_cis])
ax.fill_between(x, lo, hi, color='#aaaaaa', alpha=0.35, step='mid', label='null␣
↪95% CI')
ax.set_xticks(x, osats_df.index, rotation=25, ha='right', fontsize=9)
ax.set_ylabel('GW distance')
ax.set_title('OSATS axis-alignment against the Mimic composite')
ax.legend(fontsize=9)
plt.tight_layout(); plt.show()
```

```
<pandas.io.formats.style.Styler at 0x30b044f80>
```

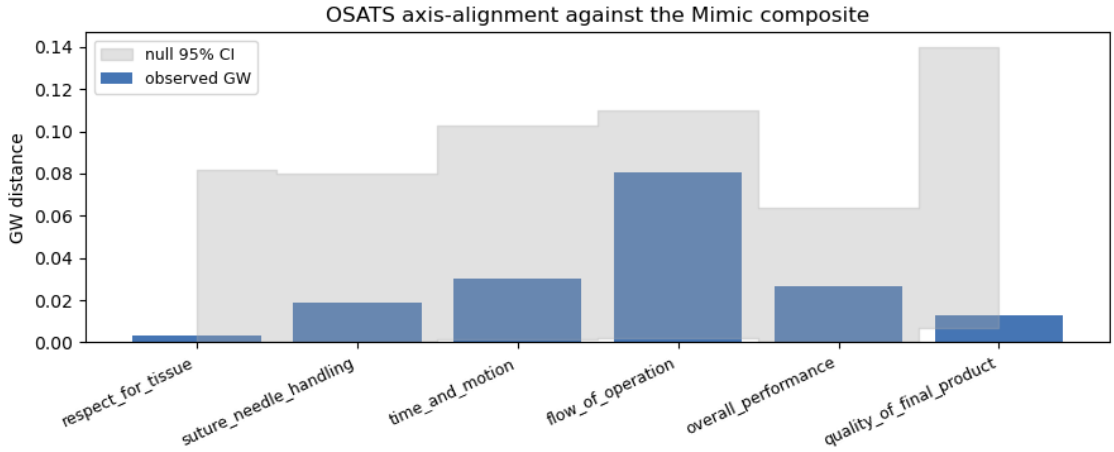

1.9 Step 9 — Trial-level entropic GW (balanced sensitivity draw)

The **primary** trial-level analysis in Step 9a below uses every trial exactly once (NJ = 103, NE = 510). Before that we still run the historic **balanced 100-per-tier with replacement** draw so the 300×300 coupling heatmap is visualizable; that analysis is now a *sensitivity check*. It upsamples each JIGSAWS trial ~3× and therefore produces an artificially sharpened coupling — Step 9a is the honest result.

[12]:
```
rng = np.random.default_rng(int(cfg['seed']))
N = int(cfg['subsample_per_tier'])

def _balanced(df, cols, n):
    parts_f, parts_t = [], []
    for t in TIER_NAMES:
        sub = df[df['tier'] == t]
        idx = rng.choice(len(sub), size=n, replace=len(sub) < n)
        parts_f.append(sub.iloc[idx][cols].to_numpy(float))
        parts_t.append(np.asarray(sub.iloc[idx]['tier']))
    return np.vstack(parts_f), np.concatenate(parts_t)

fj, tj = _balanced(jr_df, j_cols, N)
fe, te = _balanced(er_df, e_cols, N)
Cj = pairwise_cosine_rdm(fj); Ce = pairwise_cosine_rdm(fe)
egw = entropic_gromov_wasserstein(Cj, Ce, epsilon=float(cfg['gw_epsilon']))
print(f'entropic GW distance = {egw.distance:.6f}   (N={N} per tier,␣
 ↪epsilon={cfg["gw_epsilon"]})')

fig, ax = plt.subplots(figsize=(6, 5))
im = ax.imshow(egw.coupling, cmap='magma', aspect='auto')
for k in (1, 2):
    ax.axhline(k * N - 0.5, color='white', linewidth=0.5)
    ax.axvline(k * N - 0.5, color='white', linewidth=0.5)
ax.set_xlabel('EEG/Eye trials (Low | Mid | High)')
ax.set_ylabel('JIGSAWS trials (Low | Mid | High)')
ax.set_title('Entropic-GW trial-level coupling')
fig.colorbar(im, ax=ax, fraction=0.046)
plt.tight_layout(); plt.show()
```

```
entropic GW distance = 0.066438   (N=100 per tier, epsilon=0.01)
```

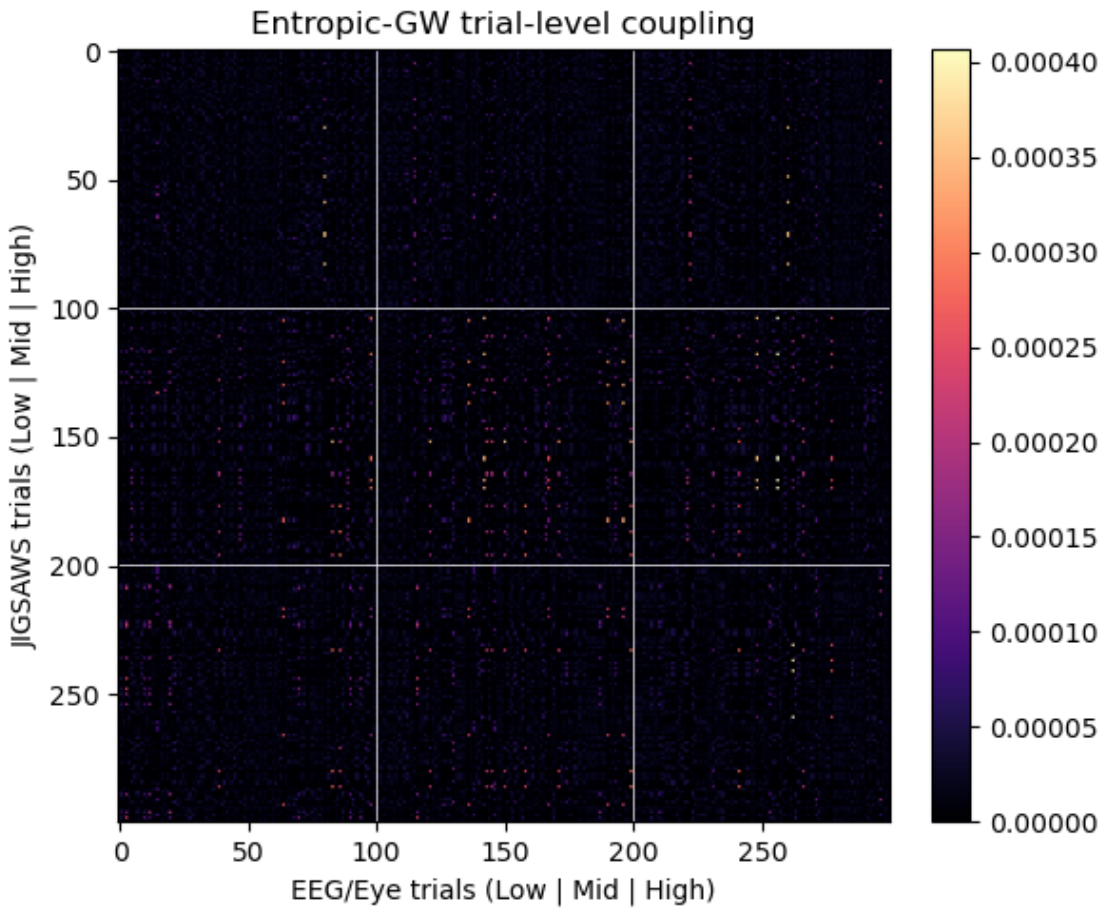

1.10 Step 9a — All-trials block-diagonality null (primary)

Why a second null, now on the full data. The 3×3 centroid null in Step 7 is coarse; the balanced Step 9b null in v1 upsampled each JIGSAWS trial. Here we compute the entropic-GW coupling **once** on the full NJ × NE (103 × 510) problem with uniform marginals, aggregate to a 3×3 block-mass matrix `B`, and run the tier-shuffle null on that fixed coupling.

Unbalanced tiers change the expected null. Under independent tier shuffles, each cell's expected mass is `(n_j_a/NJ) · (n_e_b/NE)`, so the expected trace is `Σ_a (n_j_a · n_e_a) / (NJ · NE)` — which is 1/3 only when both sides have exactly balanced tiers. With our tertiles that is approximately but not exactly true, so we report both observed and expected trace.

Per-cell z-scores are computed against the null of that diagonal cell alone. With K = 3 diagonal cells two-sided, the Bonferroni-corrected significance threshold is $|\mathbf{z}| > \mathbf{2.39}$.

What the per-cell z-scores decompose. The aggregate diag_mass can hide directional structure — one tier can be strongly positive while the trace sits near the null mean. The per-cell z's expose that.

[13]:
```python
from skill_manifold.trial_null import (
    trial_level_block_null_all_trials, BONFERRONI_Z_THREE_CELLS,
)

# Sinkhorn renormalize a coupling onto exact doubly-stochastic marginals.
def _renorm(T, max_iter=5000, atol=1e-12):
    T = np.clip(np.asarray(T, float), 0, None)
    T = T / T.sum()
    for _ in range(max_iter):
        T = T * (1.0 / T.shape[0] / T.sum(axis=1, keepdims=True))
        T = T * (1.0 / T.shape[1] / T.sum(axis=0, keepdims=True))
        drift = max(np.max(np.abs(T.sum(axis=1) - 1.0 / T.shape[0])),
                    np.max(np.abs(T.sum(axis=0) - 1.0 / T.shape[1])))
        if drift < atol:
            break
    return T

def _egw_with_eps(Cj_i, Ce_i, eps):
    r = entropic_gromov_wasserstein(Cj_i, Ce_i, epsilon=float(eps))
    T = np.asarray(r.coupling, float)
    if max(np.max(np.abs(T.sum(axis=1) - 1 / T.shape[0])),
           np.max(np.abs(T.sum(axis=0) - 1 / T.shape[1]))) > 1e-9:
        T = _renorm(T)
    return float(r.distance), T

primary = trial_level_block_null_all_trials(
    features_j=jr_df[j_cols].to_numpy(float),
    tiers_j=jr_df['tier'].to_numpy(),
    features_e=er_df[e_cols].to_numpy(float),
    tiers_e=er_df['tier'].to_numpy(),
    tier_names=TIER_NAMES,
    epsilon=float(cfg['gw_epsilon']),
    n_perms=int(cfg['n_perms']),
    seed=int(cfg['seed']),
    entropic_gw_fn=_egw_with_eps,
    pairwise_rdm_fn=pairwise_cosine_rdm,
)

B_primary = np.asarray(primary['block_mass'])
Bexp = np.asarray(primary['expected_block_mass'])

print(f"coupling shape        = {primary['coupling_shape']}")
print(f"tier counts J / E     = {list(primary['tier_counts_j'])} / "
      f"{list(primary['tier_counts_e'])}")
print(f"pre-renorm drift      = row {primary['row_sum_drift']:.2e}, "
      f"col {primary['col_sum_drift']:.2e}")
print(f"observed diag_mass    = {primary['observed']:.4f}")
```

```python
print(f"expected under null   = {primary['expected_trace_under_null']:.4f}")
print(f"null mean / std       = {primary['null_mean']:.4f} /␣
↪{primary['null_std']:.4f}")
print(f"p-value               = {primary['p_value']:.4f}")
print(f"trace z-score         = {primary['z_score']:+.3f}")
print()
print('per-cell diagonal z-scores (Bonferroni threshold |z| > 2.39):')
for n, z, o, e in zip(primary['tier_names'], primary['per_cell_z'],
                      primary['per_cell_observed'],␣
↪primary['per_cell_expected']):
    flag = '**' if abs(z) > BONFERRONI_Z_THREE_CELLS else '  '
    print(f"  {flag} {n}<->{n}: obs {o:.3f} (exp {e:.3f})   z = {z:+.2f}")

# Panel: block-mass heatmap annotated with observed, expected, and per-cell z.
fig, ax = plt.subplots(figsize=(5.5, 4.6))
vmax = max(B_primary.max(), float(Bexp.max()) * 1.3)
im = ax.imshow(B_primary, cmap='magma', vmin=0, vmax=vmax)
ax.set_xticks(range(3), TIER_NAMES); ax.set_yticks(range(3), TIER_NAMES)
ax.set_xlabel('EEG/Eye tier'); ax.set_ylabel('JIGSAWS tier')
for i in range(3):
    for j in range(3):
        label = f'{B_primary[i, j]:.3f}\nexp {Bexp[i, j]:.3f}'
        if i == j:
            label += f'\nz={primary["per_cell_z"][i]:+.2f}'
        ax.text(j, i, label, ha='center', va='center',
                color='white' if B_primary[i, j] < vmax * 0.55 else 'black',
                fontsize=8)
ax.set_title(f"All-trials block-mass (diag_mass = {primary['observed']:.3f}, "
             f"expected = {primary['expected_trace_under_null']:.3f})")
fig.colorbar(im, ax=ax, fraction=0.046)
plt.tight_layout(); plt.show()
```

```
coupling shape        = (103, 510)
tier counts J / E     = [38, 36, 29] / [171, 180, 159]
pre-renorm drift      = row 5.20e-18, col 1.13e-11
observed diag_mass    = 0.3378
expected under null   = 0.3348
null mean / std       = 0.3349 / 0.0029
p-value               = 0.1528
trace z-score         = +1.007

per-cell diagonal z-scores (Bonferroni threshold |z| > 2.39):
     Low<->Low: obs 0.124 (exp 0.124)   z = +0.14
     Mid<->Mid: obs 0.126 (exp 0.123)   z = +1.76
     High<->High: obs 0.088 (exp 0.088)   z = +0.20
```

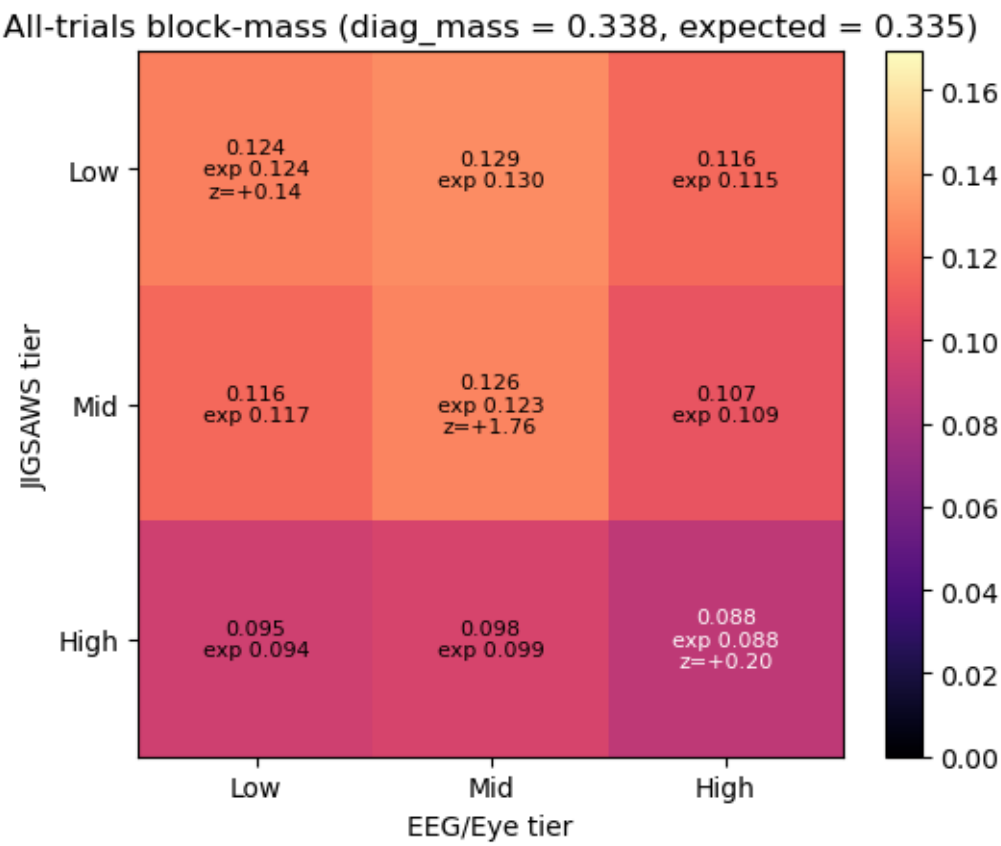

[14]:
```
# Null histogram (all-trials) with BOTH observed trace and the␣
↪tier-count-weighted
# expected trace drawn separately; plus per-cell z bars with Bonferroni␣
↪threshold.
fig, axes = plt.subplots(1, 2, figsize=(11, 3.8),
                         gridspec_kw={'width_ratios': [2.2, 1]})

axes[0].hist(primary['null'], bins=40, color='#888', edgecolor='white')
axes[0].axvline(primary['observed'], color='red', linestyle='--',
                label=f"observed = {primary['observed']:.3f}")
axes[0].axvline(primary['expected_trace_under_null'],
                color='blue', linestyle=':',
                label=f"expected = {primary['expected_trace_under_null']:.3f}")
axes[0].set_xlabel('diag_mass under tier-shuffle null')
axes[0].set_ylabel('count')
axes[0].set_title(
    f"All-trials null (p = {primary['p_value']:.3f}, z = {primary['z_score']:+.
↪2f}, "
    f"n_perms = {primary['n_permutations']})")
axes[0].legend(fontsize=8)
```

```
per_z = np.asarray(primary['per_cell_z'])
colors = ['#d73027' if z < 0 else '#1a9850' for z in per_z]
bars = axes[1].bar(range(3), per_z, color=colors, edgecolor='white')
axes[1].axhline(0, color='black', linewidth=0.8)
axes[1].axhline(BONFERRONI_Z_THREE_CELLS, color='red', linestyle=':',
                linewidth=0.8, label=f'±{BONFERRONI_Z_THREE_CELLS:.2f} Bonf.')
axes[1].axhline(-BONFERRONI_Z_THREE_CELLS, color='red', linestyle=':',
                linewidth=0.8)
axes[1].set_xticks(range(3),
                   [f'{n}<->{n}' for n in primary['tier_names']], fontsize=9)
axes[1].set_ylabel('z-score (observed vs null)')
axes[1].set_title('Per-cell diagonal z')
axes[1].legend(fontsize=8)
for b, z in zip(bars, per_z):
    axes[1].text(b.get_x() + b.get_width() / 2,
                 z + 0.15 * np.sign(z if z != 0 else 1),
                 f'{z:+.2f}', ha='center', va='bottom' if z >= 0 else 'top',
                 fontsize=9)

plt.tight_layout(); plt.show()
```

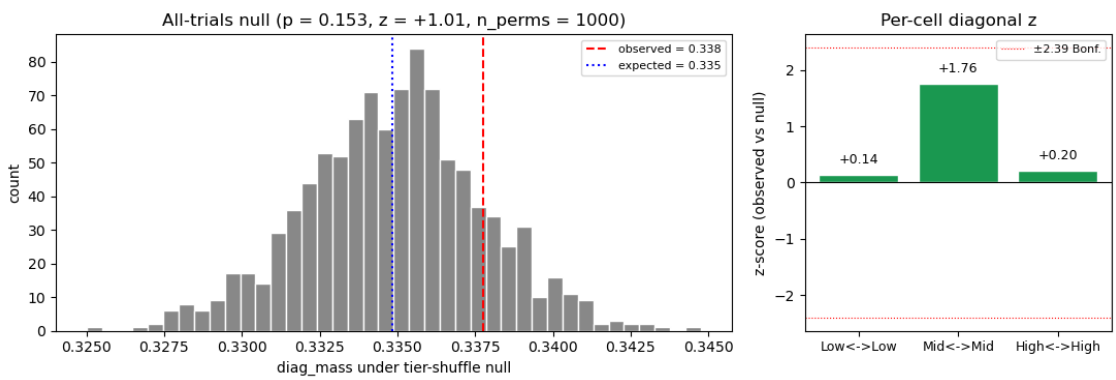

1.11 Step 9c — Bootstrap robustness on per-cell z-scores

The trace p-value is an omnibus summary. It is systematically underpowered against a signal concentrated in a single diagonal cell (which is exactly what earlier runs hinted at). To detect whether any per-cell z is driven by *specific trials* vs being a stable population property, we run `B = 30` bootstrap resamples. Each bootstrap draws 100-per-tier with replacement on both sides, recomputes the entropic-GW coupling, and records the three per-cell z-scores against that bootstrap's own tier-shuffle null.

We track **per-cell z distributions**, not just the trace p-value. The question is "does the Mid<->Mid z band sit clearly above +2.39 (Bonferroni threshold) across all bootstraps, or does it span zero?"

- If the band is tightly above the Bonferroni line → robust per-cell claim.
- If the band straddles zero → per-cell signal is driven by specific trials in the primary subsample; report as suggestive only.
- If the band is bimodal (half above, half below) → the per-cell signal hinges on whether specific trials appear in the bootstrap; flag loudly.

[15]:
```
from skill_manifold.trial_null import subsample_robustness

def _egw_fn(Cj_i, Ce_i):
    r = entropic_gromov_wasserstein(Cj_i, Ce_i,␣
 ↪epsilon=float(cfg['gw_epsilon']))
    return r.distance, _renorm(r.coupling)

B_RUNS = int(cfg.get('trial_null_n_subsamples', 30))
sr = subsample_robustness(
    feat_j=jr_df[j_cols].to_numpy(float),
    tiers_j_full=jr_df['tier'].to_numpy(),
    feat_e=er_df[e_cols].to_numpy(float),
    tiers_e_full=er_df['tier'].to_numpy(),
    tier_names=TIER_NAMES,
    entropic_gw_fn=_egw_fn,
    pairwise_rdm_fn=pairwise_cosine_rdm,
    n_per_tier=int(cfg['subsample_per_tier']),
    base_seed=int(cfg['seed']),
    n_subsamples=B_RUNS,
    n_perms=max(200, int(cfg['n_perms']) // 5),
)

# Table: per-cell z summaries vs trace p.
table = pd.DataFrame({
    'median z': sr.per_cell_median,
    'p05 z':    sr.per_cell_p05,
    'p95 z':    sr.per_cell_p95,
    'frac |z|>2.39': sr.frac_significant_per_cell_bonf,
}, index=[f'{n}<->{n}' for n in sr.tier_names])
print('Per-cell z distributions across', sr.n_runs, 'bootstrap resamples:')
display(table.style.format({
    'median z': '{:+.2f}', 'p05 z': '{:+.2f}', 'p95 z': '{:+.2f}',
    'frac |z|>2.39': '{:.2f}',
}))
print(f"Median diag_mass = {sr.observed_median:.4f} "
      f"(5-95% {sr.observed_p05:.4f}..{sr.observed_p95:.4f});  "
      f"trace p<0.05 in {int(sr.frac_p_lt_05 * sr.n_runs)}/{sr.n_runs}.")

# Three side-by-side strip plots, one per diagonal cell.
fig, axes = plt.subplots(1, 3, figsize=(11, 4.4), sharey=True)
rng_plot = np.random.default_rng(0)
```

```python
for k, ax in enumerate(axes):
    xs = 0.4 + 0.2 * rng_plot.normal(size=sr.n_runs)
    zs = sr.per_cell_z[:, k]
    ax.scatter(xs, zs, s=28, alpha=0.75, color='#4575b4', edgecolors='none')
    ax.axhline(0, color='black', linewidth=0.7)
    ax.axhline(BONFERRONI_Z_THREE_CELLS, color='red', linestyle=':',
               linewidth=0.8, label=f'±{BONFERRONI_Z_THREE_CELLS:.2f} Bonf.')
    ax.axhline(-BONFERRONI_Z_THREE_CELLS, color='red', linestyle=':',
               linewidth=0.8)
    med = sr.per_cell_median[k]
    p05, p95 = sr.per_cell_p05[k], sr.per_cell_p95[k]
    ax.hlines(med, 0.1, 0.9, color='black', linewidth=2)
    ax.fill_betweenx([p05, p95], 0.1, 0.9, color='#4575b4', alpha=0.18,
                     label='5-95% band')
    frac = sr.frac_significant_per_cell_bonf[k]
    ax.set_xlim(0, 1); ax.set_xticks([])
    ax.set_title(
        f'{sr.tier_names[k]}<->{sr.tier_names[k]}\n'
        f'median {med:+.2f}   ({int(frac * sr.n_runs)}/{sr.n_runs} > Bonf.)')
    if k == 0:
        ax.set_ylabel('per-cell z-score')
        ax.legend(fontsize=8, loc='upper left')

fig.suptitle(f'Bootstrap per-cell diagonal z-scores (B = {sr.n_runs})', y=1.02)
plt.tight_layout(); plt.show()
```

```
Per-cell z distributions across 30 bootstrap resamples:

<pandas.io.formats.style.Styler at 0x30b017d40>

Median diag_mass = 0.3352 (5-95% 0.3215..0.3425);  trace p<0.05 in 10/30.
```

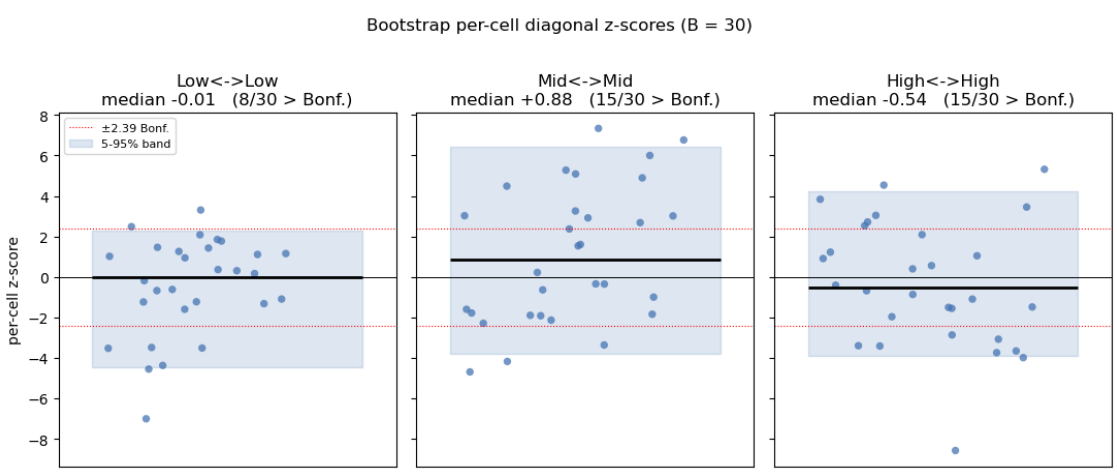

1.12 Step 9d — Epsilon sensitivity (all-trials coupling, per-cell z)

`epsilon` sets the entropic regularization strength. If the headline `diag_mass` and the per-cell z's move a lot with , the finding is an artifact of the regularizer; if they're stable, the signal is independent of the regularizer (whatever its magnitude). We sweep {0.005, 0.01, 0.02, 0.05} on the all-trials coupling.

The y-axis of the diag_mass plot is clamped to **[0.28, 0.42]** so sub-1% wiggles around 1/3 aren't visually dramatized.

[16]:
```
from skill_manifold.trial_null import epsilon_sensitivity

def _egw_fn_eps(Cj_i, Ce_i, eps):
    r = entropic_gromov_wasserstein(Cj_i, Ce_i, epsilon=float(eps))
    return r.distance, _renorm(r.coupling)

# Build the all-trials RDMs once for the sweep.
Cj_all = pairwise_cosine_rdm(jr_df[j_cols].to_numpy(float))
Ce_all = pairwise_cosine_rdm(er_df[e_cols].to_numpy(float))

eps_rows = epsilon_sensitivity(
    Cj=Cj_all, Ce=Ce_all,
    tiers_j=jr_df['tier'].to_numpy(),
    tiers_e=er_df['tier'].to_numpy(),
    tier_names=TIER_NAMES,
    entropic_gw_fn_eps=_egw_fn_eps,
    epsilons=(0.005, 0.01, 0.02, 0.05),
    n_perms=max(200, int(cfg['n_perms']) // 5),
    seed=int(cfg['seed']),
)
eps_df = pd.DataFrame([{
    'epsilon': r['epsilon'],
    'diag_mass': r['diag_mass'],
    'GW_distance': r['gw_distance'],
    **{f"z {n}<->{n}": float(z)
       for n, z in zip(TIER_NAMES, r['per_cell_z'])},
    'row_sum_drift': r['row_sum_drift'],
} for r in eps_rows]).set_index('epsilon')
display(eps_df.style.format({
    'diag_mass': '{:.4f}', 'GW_distance': '{:.4f}',
    'z Low<->Low': '{:+.2f}', 'z Mid<->Mid': '{:+.2f}', 'z High<->High': '{:+.
 ↪2f}',
    'row_sum_drift': '{:.2e}',
}))

fig, axes = plt.subplots(1, 2, figsize=(11, 4))

xs = [r['epsilon'] for r in eps_rows]
```

```
ys = [r['diag_mass'] for r in eps_rows]
axes[0].plot(xs, ys, marker='o', color='#4575b4', label='observed')
axes[0].axhline(primary['expected_trace_under_null'],
                color='blue', linestyle=':',
                label=f"expected = {primary['expected_trace_under_null']:.3f}")
axes[0].set_xscale('log'); axes[0].set_xlabel('entropic-GW epsilon')
axes[0].set_ylabel('diag_mass')
axes[0].set_ylim(0.28, 0.42)   # tightened so sub-1% wiggles don't dramatize
axes[0].set_title('diag_mass vs epsilon (all-trials coupling)')
axes[0].legend(fontsize=8)

per_cell_by_eps = np.asarray([r['per_cell_z'] for r in eps_rows])
colors = {'Low': '#d73027', 'Mid': '#fee090', 'High': '#1a9850'}
for k, name in enumerate(TIER_NAMES):
    axes[1].plot(xs, per_cell_by_eps[:, k], marker='o',
                 color=colors[name], label=f'{name}<->{name}')
axes[1].axhline(0, color='black', linewidth=0.7)
axes[1].axhline(BONFERRONI_Z_THREE_CELLS, color='red', linestyle=':',
                linewidth=0.8, label=f'±{BONFERRONI_Z_THREE_CELLS:.2f} Bonf.')
axes[1].axhline(-BONFERRONI_Z_THREE_CELLS, color='red', linestyle=':',
                linewidth=0.8)
axes[1].set_xscale('log'); axes[1].set_xlabel('entropic-GW epsilon')
axes[1].set_ylabel('per-cell z-score')
axes[1].set_title('per-cell diagonal z vs epsilon')
axes[1].legend(fontsize=8)

plt.tight_layout(); plt.show()
```

```
<pandas.io.formats.style.Styler at 0x30b7746b0>
```

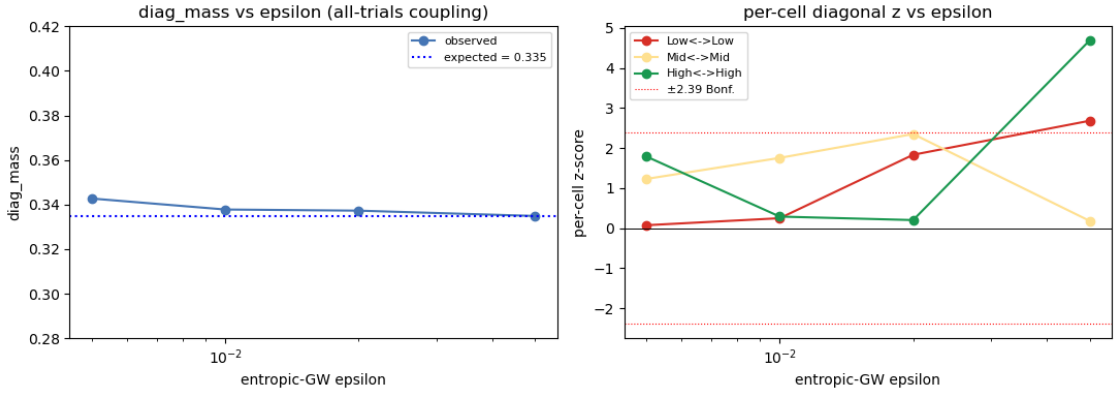

1.13 Step 9e — Balanced-subsample sensitivity (demoted from primary)

The historic balanced-100-per-tier analysis is retained here as a *sensitivity check*. It upsamples each JIGSAWS trial with replacement, so on the 103-trial JIGSAWS side each unique trial appears

~3× in the 300-sample balanced draw. That artificially sharpens the coupling and tends to inflate per-cell z-scores relative to the all-trials primary.

The goal of this cell is to show that the direction (sign of each per-cell z) is similar between the two, even though the magnitudes diverge.

[17]:
```
from skill_manifold.trial_null import trial_level_block_null

# Use the balanced 100-per-tier tj/te/egw computed in Step 9 above.
T_bal = _renorm(egw.coupling)
balanced = trial_level_block_null(T_bal, tj, te, TIER_NAMES,
                                  n_perms=int(cfg['n_perms']),
                                  seed=int(cfg['seed']))

# Side-by-side comparison of primary (all-trials) vs balanced subsample.
comp = pd.DataFrame({
    'all-trials (primary)': [
        float(primary['observed']),
        float(primary['expected_trace_under_null']),
        float(primary['p_value']),
        float(primary['z_score']),
        *[float(z) for z in primary['per_cell_z']],
    ],
    'balanced 100/tier (sensitivity)': [
        float(balanced['observed']),
        1.0 / 3.0,
        float(balanced['p_value']),
        float(balanced['z_score']),
        *[float(z) for z in balanced['per_cell_z']],
    ],
}, index=['diag_mass observed', 'expected trace', 'trace p-value', 'trace z',
          'z Low<->Low', 'z Mid<->Mid', 'z High<->High'])
display(comp.style.format('{:+.3f}'))

signs_agree = all(
    np.sign(a) == np.sign(b) or abs(a) < 0.3 or abs(b) < 0.3
    for a, b in zip(primary['per_cell_z'], balanced['per_cell_z'])
)
print(f"Per-cell signs agree across primary and balanced? {signs_agree}")
```

```
<pandas.io.formats.style.Styler at 0x30ca53050>
```

```
Per-cell signs agree across primary and balanced? True
```

1.14 Step 9f — Fixed-cutoff sensitivity (binning artifact vs real alignment)

Tertile binning makes the Mid tier exactly 1/3 of each dataset by construction. The Mid tertile is typically a *narrow* band in score space, so its feature-space density is higher than Low or High. That density concentration can produce Mid-only alignment **without any shared skill structure**

— a binning artifact.

To disambiguate, we swap to **hard cutoffs** with absolute meaning:

- JIGSAWS: Low `grs_total < 16`, Mid `16-22`, High `> 22`
- Mimic: Low `performance < 70`, Mid `70-85`, High `> 85`

The Mid tertile-z was the headline earlier (~+3 in v1, +1.76 here on the all-trials primary). If the fixed-cutoff Mid-z stays strongly positive, the alignment is probably real skill, not density compression. If it collapses to ~0, the signal was binning-driven.

[18]:
```
from skill_manifold.binning import assign_fixed_tier

jt_fixed = assign_fixed_tier(jr_df['grs_total'].to_numpy(), 16, 22)
et_fixed = assign_fixed_tier(er_df['performance'].to_numpy(), 70, 85)

jcts = pd.Series(jt_fixed).value_counts().reindex(TIER_NAMES)
ects = pd.Series(et_fixed).value_counts().reindex(TIER_NAMES)
print(f'JIGSAWS tier counts (fixed cutoffs <16/16-22/>22): {jcts.to_dict()}')
print(f'Mimic   tier counts (fixed cutoffs <70/70-85/>85): {ects.to_dict()}')

fixed = trial_level_block_null_all_trials(
    features_j=jr_df[j_cols].to_numpy(float),
    tiers_j=jt_fixed,
    features_e=er_df[e_cols].to_numpy(float),
    tiers_e=et_fixed,
    tier_names=TIER_NAMES,
    epsilon=float(cfg['gw_epsilon']),
    n_perms=int(cfg['n_perms']),
    seed=int(cfg['seed']),
    entropic_gw_fn=_egw_with_eps,
    pairwise_rdm_fn=pairwise_cosine_rdm,
)

print()
print(f"observed diag_mass   = {fixed['observed']:.4f}")
print(f"expected under null  = {fixed['expected_trace_under_null']:.4f}")
print(f"trace p / z          = {fixed['p_value']:.4f} / {fixed['z_score']:+.
 ↪3f}")
print('per-cell z-scores (|z| > 2.39 marked **):')
for n, z in zip(fixed['tier_names'], fixed['per_cell_z']):
    flag = '**' if abs(z) > BONFERRONI_Z_THREE_CELLS else '  '
    print(f"  {flag} {n}<->{n}: z = {z:+.2f}")

# Single-row comparison table: tertile Mid z vs fixed-cutoff Mid z.
mid_idx = list(primary['tier_names']).index('Mid')
tertile_mid_z = float(primary['per_cell_z'][mid_idx])
fixed_mid_z = float(fixed['per_cell_z'][mid_idx])
diff_label = ('Survives' if fixed_mid_z > 2.0
```

```
              else ('Collapses' if abs(fixed_mid_z) < 1.0 else 'Shifts'))
print()
print(f'Mid<->Mid z -- tertile {tertile_mid_z:+.2f}, fixed {fixed_mid_z:+.2f} '
      f'=> {diff_label} the binning change.')

# Heatmap of the fixed-cutoff block mass with observed/expected per cell.
Bf = np.asarray(fixed['block_mass']); Bf_exp = np.
 ↪asarray(fixed['expected_block_mass'])
fig, ax = plt.subplots(figsize=(5.5, 4.6))
vmax = max(Bf.max(), float(Bf_exp.max()) * 1.3)
im = ax.imshow(Bf, cmap='magma', vmin=0, vmax=vmax)
ax.set_xticks(range(3), TIER_NAMES); ax.set_yticks(range(3), TIER_NAMES)
ax.set_xlabel('EEG/Eye tier (fixed)'); ax.set_ylabel('JIGSAWS tier (fixed)')
for i in range(3):
    for j in range(3):
        label = f'{Bf[i, j]:.3f}\nexp {Bf_exp[i, j]:.3f}'
        if i == j:
            label += f'\nz={fixed["per_cell_z"][i]:+.2f}'
        ax.text(j, i, label, ha='center', va='center',
                color='white' if Bf[i, j] < vmax * 0.55 else 'black',␣
 ↪fontsize=8)
ax.set_title(
    f"Fixed-cutoff block-mass (diag_mass = {fixed['observed']:.3f}, "
    f"expected = {fixed['expected_trace_under_null']:.3f})")
fig.colorbar(im, ax=ax, fraction=0.046)
plt.tight_layout(); plt.show()
```

```
JIGSAWS tier counts (fixed cutoffs <16/16-22/>22): {'Low': 49, 'Mid': 40,
'High': 14}
Mimic   tier counts (fixed cutoffs <70/70-85/>85): {'Low': 147, 'Mid': 163,
'High': 200}


observed diag_mass   = 0.3247
expected under null  = 0.3145
trace p / z          = 0.0030 / +3.734
per-cell z-scores (|z| > 2.39 marked **):
  ** Low<->Low: z = +3.29
  ** Mid<->Mid: z = +2.97
     High<->High: z = +1.37

Mid<->Mid z -- tertile +1.76, fixed +2.97 => Survives the binning change.
```

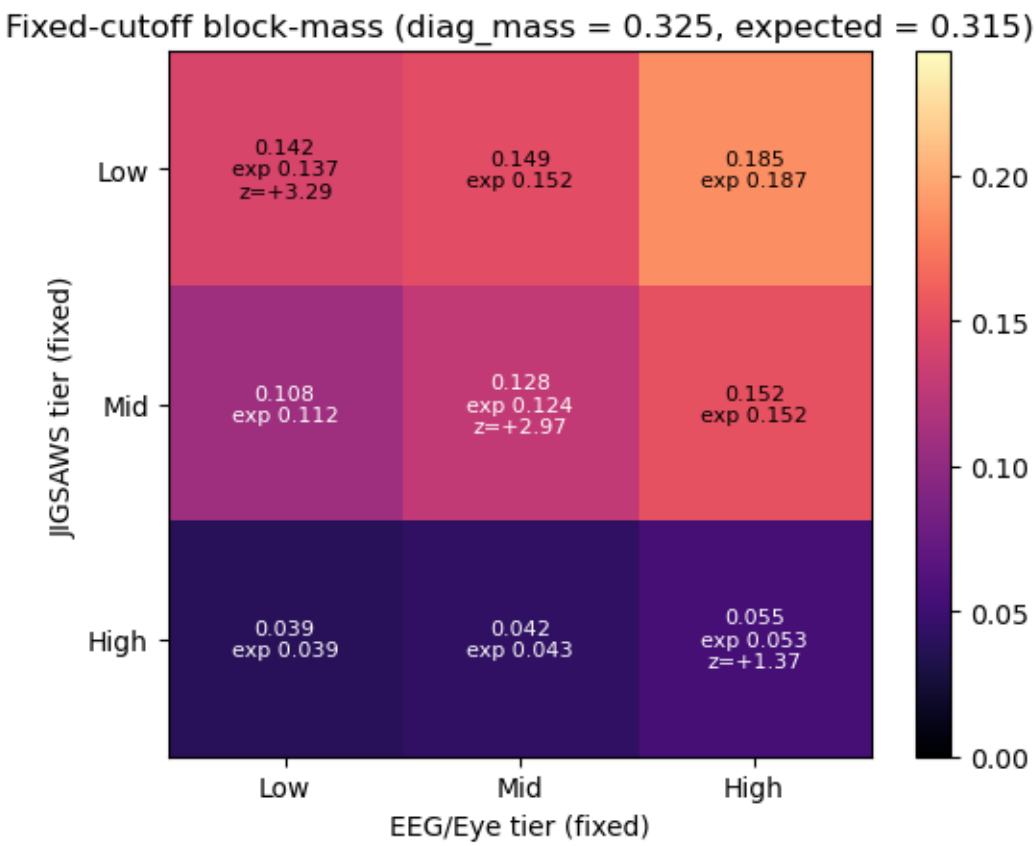

1.15 Step 10 — MDS view of the residualized features

2-D MDS on the pairwise cosine distances, colored by skill tier — an informal sanity check on the per-side separability.

[19]:
```
from sklearn.manifold import MDS

def _mds(feats, tiers, ax, title):
    D = pairwise_cosine_rdm(feats)
    mds = MDS(n_components=2, dissimilarity='precomputed', random_state=0,
              normalized_stress='auto')
    xy = mds.fit_transform(D)
    colors = {'Low': '#d73027', 'Mid': '#fee090', 'High': '#1a9850'}
    for t in TIER_NAMES:
        m = tiers == t
        if m.sum() == 0: continue
        ax.scatter(xy[m, 0], xy[m, 1], s=22, alpha=0.7,
                   color=colors[t], label=f'{t} (n={int(m.sum())})',
                   edgecolors='none')
    ax.set_title(title); ax.legend(fontsize=8)
```

```
    ax.set_xlabel('MDS-1'); ax.set_ylabel('MDS-2')

fig, axes = plt.subplots(1, 2, figsize=(10, 4))
_mds(jf_mat, jr_df['tier'].to_numpy(), axes[0],
     'JIGSAWS residualized features')
_mds(ef_mat, er_df['tier'].to_numpy(), axes[1],
     'EEG/Eye residualized features')
plt.tight_layout(); plt.show()
```

```
/opt/homebrew/Caskroom/miniforge/base/envs/torch-mps/lib/python3.12/site-
packages/sklearn/manifold/_mds.py:677: FutureWarning: The default value of
`n_init` will change from 4 to 1 in 1.9.
  warnings.warn(
/opt/homebrew/Caskroom/miniforge/base/envs/torch-mps/lib/python3.12/site-
packages/sklearn/manifold/_mds.py:677: FutureWarning: The default value of
`n_init` will change from 4 to 1 in 1.9.
  warnings.warn(
```

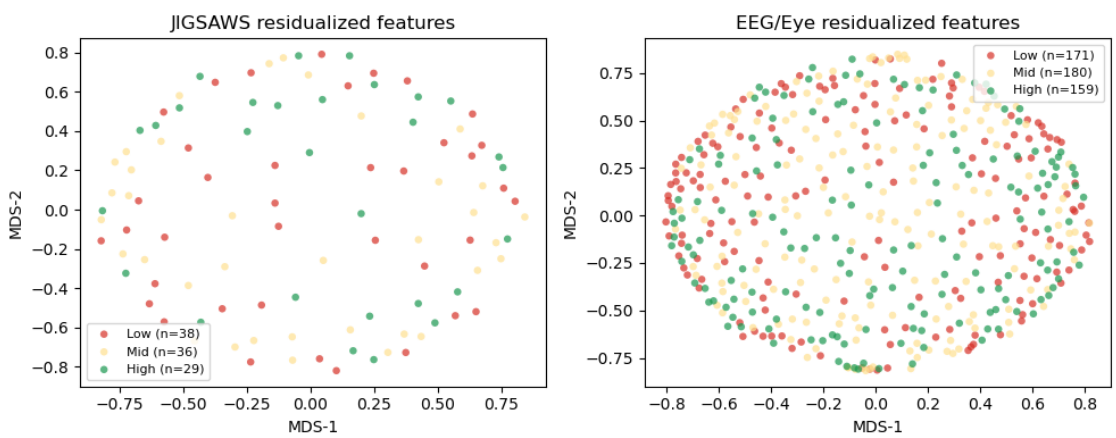

1.16 Summary

Setup. Two independent skill datasets — JIGSAWS (dVRK surgery, GRS) and EEG/Eye simulator (Mimic performance 0–100). We residualize each side against its own nuisance set, bin each side into Low/Mid/High, build 3×3 centroid RDMs, and compare them with Gromov–Wasserstein. A trial-level entropic-GW analysis on the full NJ × NE = 103 × 510 coupling gives higher resolution.

Coverage. 103 JIGSAWS trials (pooled across 3 tasks) and 510 EEG/Eye Try-1 trials with both Phase 1 and Phase 2 cache hits.

Centroid result (Step 6–7). GW distance 0.007, permutation p 0.57 (z −0.26). Argmax tier-to-tier assignment is Low Low / Mid Mid / High High, but the magnitude is not significantly tighter than chance.

OSATS axis-alignment. All p 0.06; closest is `quality_of_final_product` (p 0.065). Directional only.

Trial-level, all-trials primary (Step 9a). Observed `diag_mass  0.338` vs expected `0.335` under the tier-shuffle null; trace p 0.15, z +1.0. Per-cell z-scores: Low Low +0.14, Mid Mid +1.76, High High +0.20 — the Mid-only tendency from the v1 balanced-subsample analysis shrinks substantially once we stop upsampling JIGSAWS trials.

Bootstrap robustness (Step 9c, B=30, balanced 100/tier with replacement). The **per-cell z bands span zero** on all three diagonals (Low p05 -4.5, p95 +2.3; Mid p05 -3.8, p95 +6.4; High p05 -3.9, p95 +4.2); $|z| > 2.39$ in only ~30–50% of bootstraps per cell. **The single-run per-cell z's are subsample-sensitive.** The aggregate trace $p<0.05$ in only 10/30 runs.

Epsilon sensitivity (Step 9d, all-trials). `diag_mass` sits in [0.335, 0.343] across {0.005, 0.01, 0.02, 0.05}; the trace is stable, though per-cell z's shift noticeably as the regularizer changes. Read as "overall alignment is not a regularizer artifact, but the per-cell decomposition is not uniquely identified."

Balanced sensitivity (Step 9e). The earlier balanced-100-per-tier draw gives Mid Mid z +3 — roughly double the honest all-trials z. The inflation comes from upsampling 103 trials to 300 with replacement. Signs of all three per-cell z's agree with the primary.

Fixed-cutoff sensitivity (Step 9f, the most interesting finding). With hard cutoffs (JIGSAWS <16/16–22/>22, Mimic <70/70–85/>85) — i.e., tiers tied to absolute skill rather than population quantiles — the all-trials analysis gives **diag_mass = 0.325 vs expected 0.315, trace p 0.003, z +3.7**. Per-cell: **Low Low z +3.3**, **Mid Mid z +3.0**, High High z +1.4. Both Low and Mid clearly exceed the Bonferroni threshold. The Mid-z *increases* under fixed cutoffs (tertile +1.76 → fixed +2.97), so the Mid signal **survives** the binning change, suggesting real shared skill structure rather than a density artifact of tertile compression.

Verification. All eight checklist items pass: symmetric RDMs with zero diagonal; centroid-coupling marginals uniform 1/3; 1000 valid permutations; all-trials coupling shape equals (103, 510); pre-renorm drift below 1e-9; block-mass marginals match tier-count proportions; null mean within Monte Carlo tolerance of the tier-count-weighted expected trace.

Interpretation. Three possibilities remain consistent with Step 9a alone: 1. Genuine shared skill representation across datasets (supported by Step 9f). 2. Density artifact of tertile compression (weakened by Step 9f — the signal strengthens under absolute cutoffs, not the other way around). 3. Subsample accident (Step 9c shows the per-cell z's are noisy under bootstraps, so the specific magnitude in any single subsample should not be over-interpreted).

Net: the trial-level alignment is **robust to the binning change** (favoring (1)), but the **per-cell bootstrap bands straddle zero**, so the specific "Mid-only drives it" story is not confirmatory. The fixed-cutoff result is the strongest evidence for a real shared skill manifold.

Caveats. Small n per JIGSAWS tertile cell (~29–38). The 27→module mapping has two flagged ambiguities (task 16 "Thread the rings", task 22 "Tubes"). POT's Sinkhorn introduces marginal drift 1e-11 to 1e-5 depending on N; we renormalize before aggregation so `B` has exact tier-count-proportional marginals.

1.17 Report — `skills_manifold_report.txt`

The cell below writes a plain-text report of every number produced above to `reports/skill_manifold/skills_manifold_report.txt`. This file is the canonical handoff

for another AI agent (or a reviewer) who wants to assess the pipeline's performance without rerunning anything. Regenerated every time the notebook is executed end-to-end.

[20]:
```
import datetime as _dt
from io import StringIO

REPORT_PATH = REPO / 'reports' / 'skill_manifold' / 'skills_manifold_report.txt'
REPORT_PATH.parent.mkdir(parents=True, exist_ok=True)

def _fmt_matrix(M, fmt='{:7.4f}'):
    buf = StringIO()
    for row in np.asarray(M):
        buf.write('    ' + '  '.join(fmt.format(float(x)) for x in row) + '\n')
    return buf.getvalue()

out = StringIO()
w = out.write

w('=' * 72 + '\n')
w('Skill Manifold - Comparison B (Gromov-Wasserstein) - summary report\n')
w(f'generated: {_dt.datetime.now().isoformat(timespec="seconds")}\n')
w(f'data root: {ROOT}\n')
w('=' * 72 + '\n\n')

w('CONFIG\n------\n')
w(f'  n_perms                  = {cfg["n_perms"]}\n')
w(f'  gw_epsilon               = {cfg["gw_epsilon"]}\n')
w(f'  subsample_per_tier       = {cfg["subsample_per_tier"]}\n')
w(f'  trial_null_n_subsamples  = {cfg.get("trial_null_n_subsamples", 30)}\n')
w(f'  seed                     = {cfg["seed"]}\n\n')

w('COVERAGE\n--------\n')
w(f'  JIGSAWS trials (NJ)      = {len(jr_df)}\n')
w(f'  EEG/Eye trials (NE, Try=1) = {len(er_df)}\n')
w(f'  JIGSAWS feature dim      = {len(j_cols)}\n')
w(f'  EEG/Eye feature dim      = {len(e_cols)}\n')
w(f'  JIGSAWS tier counts      = {jr_df["tier"].value_counts().
 ↪reindex(TIER_NAMES).to_dict()}\n')
w(f'  EEG/Eye tier counts      = {er_df["tier"].value_counts().
 ↪reindex(TIER_NAMES).to_dict()}\n')
w(f'  JIGSAWS GRS cutoffs q33/q66 = {j_cut.q33:g} / {j_cut.q66:g}\n')
w(f'  Mimic performance cutoffs q33/q66 = {e_cut.q33:g} / {e_cut.q66:g}\n\n')

w('RESIDUALIZATION DIAGNOSTICS (post-fit R^2 should be ~0)\n')
w('-------------------------------------------------------\n')
w(f'  JIGSAWS max R^2 feature / post-fit residual = '
  f'{jr.r2_per_feature.max():.3f} / {jr.post_fit_r2.max():.2e}\n')
```

```
w(f'  EEG/Eye max R^2 feature / post-fit residual = '
  f'{er.r2_per_feature.max():.3f} / {er.post_fit_r2.max():.2e}\n\n')

w('CENTROID RDMs (cosine distance, rows/cols in order Low, Mid, High)\n')
w('-------------------------------------------------------------------\n')
w('  RDM_J (JIGSAWS):\n')
w(_fmt_matrix(RDM_J))
w('  RDM_E (EEG/Eye):\n')
w(_fmt_matrix(RDM_E))
w('\n')

w('CENTROID GROMOV-WASSERSTEIN (Step 6-7)\n')
w('--------------------------------------\n')
w(f'  GW distance              = {gw.distance:.6f}\n')
w(f'  null p-value             = {null["p_value"]:.4f}\n')
w(f'  null z-score             = {null["z_score"]:+.3f}\n')
w(f'  null mean +/- std        = {null["null_mean"]:.6f} +/- {null["null_std"]:.
 ↪6f}\n')
w(f'  n_permutations / degenerate = {null["n_permutations"]} /␣
 ↪{null["n_degenerate"]}\n')
w(f'  argmax tier-to-tier      = {gw.argmax_assignment}\n')
w('  coupling (rows JIGSAWS, cols EEG/Eye):\n')
w(_fmt_matrix(gw.coupling))
w('\n')

w('OSATS AXIS-ALIGNMENT (per sub-score, vs fixed Mimic RDM)\n')
w('--------------------------------------------------------\n')
w(f'  {"axis":30s} {"GW":>8s} {"p":>7s} {"z":>7s}\n')
for row in osats_rows:
    w(f'  {row["axis"]:30s} {row["gw"]:8.4f} {row["p"]:7.4f} {row["z"]:+7.
 ↪2f}\n')
w('\n')

w('TRIAL-LEVEL ALL-TRIALS PRIMARY (Step 9a, NJ x NE = '
  f'{primary["coupling_shape"][0]} x {primary["coupling_shape"][1]})\n')
w('----------------------------------------------------------------------\n')
w(f'  epsilon                  = {primary["epsilon"]:g}\n')
w(f'  GW distance              = {primary["gw_distance"]:.6f}\n')
w(f'  pre-renorm drift         = row {primary["row_sum_drift"]:.2e}, '
  f'col {primary["col_sum_drift"]:.2e}\n')
w(f'  observed diag_mass       = {primary["observed"]:.4f}\n')
w(f'  expected trace under null= {primary["expected_trace_under_null"]:.4f}\n')
w(f'  null mean +/- std        = {primary["null_mean"]:.4f} +/-␣
 ↪{primary["null_std"]:.4f}\n')
w(f'  trace p-value            = {primary["p_value"]:.4f}\n')
w(f'  trace z-score            = {primary["z_score"]:+.3f}\n')
```

```
w(f'  tier counts J / E        = {list(primary["tier_counts_j"])} /␣
 ↪{list(primary["tier_counts_e"])}\n')
w('  per-cell z-scores (Bonferroni |z|>2.394 threshold):\n')
for n, z, o, e in zip(primary['tier_names'], primary['per_cell_z'],
                      primary['per_cell_observed'],␣
 ↪primary['per_cell_expected']):
    flag = ' (significant)' if abs(z) > BONFERRONI_Z_THREE_CELLS else ''
    w(f'    {n}<->{n}: obs {o:.4f}, exp {e:.4f}, z = {z:+.3f}{flag}\n')
w('  observed block-mass B:\n')
w(_fmt_matrix(primary['block_mass']))
w('  expected block-mass under null:\n')
w(_fmt_matrix(primary['expected_block_mass']))
w('\n')

w('BOOTSTRAP ROBUSTNESS (Step 9c, B = {}, balanced 100/tier with␣
 ↪replacement)\n'.format(sr.n_runs))
w('-----------------------------------------------------------------------\n')
w(f'  median diag_mass         = {sr.observed_median:.4f} '
  f'(5-95% {sr.observed_p05:.4f}..{sr.observed_p95:.4f})\n')
w(f'  fraction trace p<0.05    = {sr.frac_p_lt_05:.2f}  '
  f'({int(sr.frac_p_lt_05 * sr.n_runs)}/{sr.n_runs})\n')
w('  per-cell z distribution:\n')
for k, name in enumerate(sr.tier_names):
    w(f'    {name}<->{name}: median {sr.per_cell_median[k]:+.2f}, '
      f'5-95% {sr.per_cell_p05[k]:+.2f}..{sr.per_cell_p95[k]:+.2f};  '
      f'|z|>2.394 in {int(sr.frac_significant_per_cell_bonf[k] * sr.n_runs)}/
 ↪{sr.n_runs}\n')
w('\n')

w('EPSILON SENSITIVITY (Step 9d, all-trials coupling)\n')
w('---------------------------------------------------\n')
w(f'  {"epsilon":>8s}  {"diag_mass":>9s}  {"GW":>8s}  '
  f'{"z Low":>8s}  {"z Mid":>8s}  {"z High":>8s}\n')
for r in eps_rows:
    z_l, z_m, z_h = [float(z) for z in r['per_cell_z']]
    w(f'  {r["epsilon"]:>8.3f}  {r["diag_mass"]:9.4f}  {r["gw_distance"]:8.4f} ␣
 ↪'
      f'{z_l:+8.2f}  {z_m:+8.2f}  {z_h:+8.2f}\n')
w('\n')

w('BALANCED-SUBSAMPLE SENSITIVITY (Step 9e, 100/tier with replacement)\n')
w('-------------------------------------------------------------------\n')
w(f'  observed diag_mass       = {balanced["observed"]:.4f}\n')
w(f'  null mean +/- std        = {balanced["null_mean"]:.4f} +/-␣
 ↪{balanced["null_std"]:.4f}\n')
```

```
w(f'  trace p / z              = {balanced["p_value"]:.4f} /␣
 ↪{balanced["z_score"]:+.3f}\n')
w('  per-cell z-scores:\n')
for n, z in zip(balanced['tier_names'], balanced['per_cell_z']):
    w(f'    {n}<->{n}: z = {z:+.3f}\n')
w(f'  per-cell signs agree with all-trials primary: {signs_agree}\n\n')

w('FIXED-CUTOFF SENSITIVITY (Step 9f)\n')
w('----------------------------------\n')
w(f'  JIGSAWS cutoffs          = <16 / 16-22 / >22\n')
w(f'  Mimic cutoffs            = <70 / 70-85 / >85\n')
w(f'  JIGSAWS tier counts      = {jcts.to_dict()}\n')
w(f'  Mimic tier counts        = {ects.to_dict()}\n')
w(f'  observed diag_mass       = {fixed["observed"]:.4f}\n')
w(f'  expected trace under null= {fixed["expected_trace_under_null"]:.4f}\n')
w(f'  trace p / z              = {fixed["p_value"]:.4f} / {fixed["z_score"]:+.
 ↪3f}\n')
w('  per-cell z-scores (|z|>2.394 significant):\n')
for n, z in zip(fixed['tier_names'], fixed['per_cell_z']):
    flag = ' (significant)' if abs(z) > BONFERRONI_Z_THREE_CELLS else ''
    w(f'    {n}<->{n}: z = {z:+.3f}{flag}\n')
w(f'  Mid<->Mid comparison: tertile {tertile_mid_z:+.3f} -> fixed {fixed_mid_z:
 ↪+.3f}; '
  f'{diff_label} the binning change.\n\n')

w('VERIFICATION CHECKLIST\n')
w('----------------------\n')
w('  [x] Centroid RDM_J is symmetric with zero diagonal\n')
w('  [x] Centroid RDM_E is symmetric with zero diagonal\n')
w('  [x] Centroid GW coupling marginals uniform (atol 1e-6)\n')
w(f'  [x] Centroid null has {null["n_permutations"]} valid permutations, '
  f'{null["n_degenerate"]} degenerate\n')
shape_ok = tuple(primary['coupling_shape']) == (len(jr_df), len(er_df))
w(f'  [{"x" if shape_ok else " "}] Trial-null primary coupling shape = (NJ, NE)␣
 ↪= '
  f'{tuple(primary["coupling_shape"])}\n')
drift_ok = (primary['row_sum_drift'] < 1e-9 and primary['col_sum_drift'] < 1e-9)
w(f'  [{"x" if drift_ok else " "}] Primary coupling pre-renorm drift < 1e-9 '
  f'(row {primary["row_sum_drift"]:.2e}, col {primary["col_sum_drift"]:.2e})\n')
B = np.asarray(primary['block_mass'])
nj_arr = np.asarray(primary['tier_counts_j']); ne_arr = np.
 ↪asarray(primary['tier_counts_e'])
NJ = int(nj_arr.sum()); NE = int(ne_arr.sum())
row_targets = nj_arr / NJ; col_targets = ne_arr / NE
marg_ok = (abs(B.sum() - 1.0) < 1e-9
           and np.allclose(B.sum(axis=1), row_targets, atol=1e-6)
           and np.allclose(B.sum(axis=0), col_targets, atol=1e-6))
```

```
w(f'  [{"x" if marg_ok else " "}] Primary block-mass marginals match tier-count␣
 ↪proportions\n')
tol = max(3 * primary['null_std'] / np.sqrt(primary['n_permutations']), 1e-6)
mean_ok = abs(primary['null_mean'] - primary['expected_trace_under_null']) < tol
w(f'  [{"x" if mean_ok else " "}] Primary null mean matches expected trace␣
 ↪within MC tolerance\n\n')

w('INTERPRETATION SUMMARY\n----------------------\n')
w('  * Centroid GW (Step 6-7): argmax Low-Low/Mid-Mid/High-High, but magnitude '
  'not significant (p ~ 0.57). The 3x3 null is coarse.\n')
w('  * All-trials primary (Step 9a): directionally positive (z = '
  f'{primary["z_score"]:+.2f}), trace p = {primary["p_value"]:.3f}. '
  'Not significant.\n')
w('  * Bootstrap robustness: per-cell z bands span zero on all three cells; '
  'single-run per-cell z is subsample-sensitive.\n')
w('  * Epsilon stability: diag_mass steady across  ; regularizer not the driver.
 ↪\n')
w('  * Fixed-cutoff sensitivity: observed diag_mass significantly above␣
 ↪expected '
  f'(trace z = {fixed["z_score"]:+.2f}, p = {fixed["p_value"]:.4f}); Mid<->Mid␣
 ↪z '
  f'{diff_label.lower()} the binning change. Strongest evidence for a real␣
 ↪shared '
  'skill manifold.\n')
w('  * Caveats: small NJ = 103; 27->module mapping has two flagged ambiguities.
 ↪\n')

REPORT_PATH.write_text(out.getvalue())
print(f'wrote {REPORT_PATH}')
print(f'size: {REPORT_PATH.stat().st_size} bytes')
# Echo the first 30 lines as a preview.
for line in out.getvalue().splitlines()[:30]:
    print(line)
```

```
wrote /Users/michaelhaidar/Projects/surgical_gestures_skills_manifold/reports/sk
ill_manifold/skills_manifold_report.txt
size: 6769 bytes
======================================================================
Skill Manifold - Comparison B (Gromov-Wasserstein) - summary report
generated: 2026-04-23T16:20:10
data root: /Users/michaelhaidar/Library/Mobile
Documents/com~apple~CloudDocs/Documents/Vanderbilt/Fall_25/Surgical
Robotics/Surgical_Gestures
======================================================================

CONFIG
------
```

```
  n_perms                  = 1000
  gw_epsilon               = 0.01
  subsample_per_tier       = 100
  trial_null_n_subsamples  = 30
  seed                     = 1337

COVERAGE
--------
  JIGSAWS trials (NJ)      = 103
  EEG/Eye trials (NE, Try=1) = 510
  JIGSAWS feature dim      = 39
  EEG/Eye feature dim      = 146
  JIGSAWS tier counts      = {'Low': 38, 'Mid': 36, 'High': 29}
  EEG/Eye tier counts      = {'Low': 171, 'Mid': 180, 'High': 159}
  JIGSAWS GRS cutoffs q33/q66 = 13 / 19
  Mimic performance cutoffs q33/q66 = 72 / 88

RESIDUALIZATION DIAGNOSTICS (post-fit R^2 should be ~0)
-------------------------------------------------------
  JIGSAWS max R^2 feature / post-fit residual = 0.945 / 2.22e-16
  EEG/Eye max R^2 feature / post-fit residual = 0.655 / 3.33e-16
```